


# Analyzing Performance Characteristics of PostgreSQL and MariaDB on NVMeVirt


2022-23583 최유진, 2022-29677 한주희

# Introduction


- I/O bandwidth limit influences both database engines, but PostgreSQL is much more sensitive than MariaDB.  
→ NVMeVirt paper concluded that *PostgreSQL is more promising on modern storage devices, whereas MariaDB is more efficient when the storage is slow.*
- *Problem* is that the paper does not provide a clear explanation of why there exists such differences.
- *Our goal* is to analyze the result with a focus on what features of database engine internals make such differences.



(c) TPS vs. target bandwidth



(e) Queue depth (MariaDB)




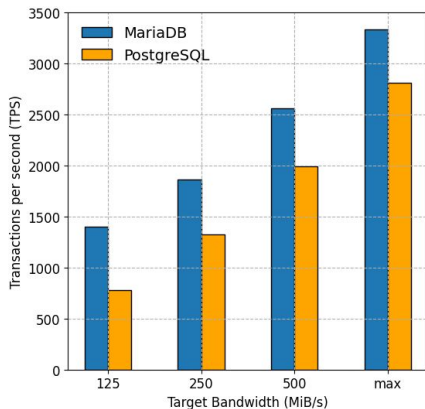
(f) Queue depth (PostgreSQL)

# Evaluation

- OLTP workload with *sysbench*
- Our maximum bandwidth is around 660 MiB/s and our observations are
  1. When the bandwidth limit is low, MariaDB outperforms PostgreSQL.
  2. PostgreSQL is more sensitive to I/O bandwidth than MariaDB.
    - Performance gap becomes smaller (1.80x to 1.19x)
  3. PostgreSQL requests higher number of pages (= higher number of I/O requests).

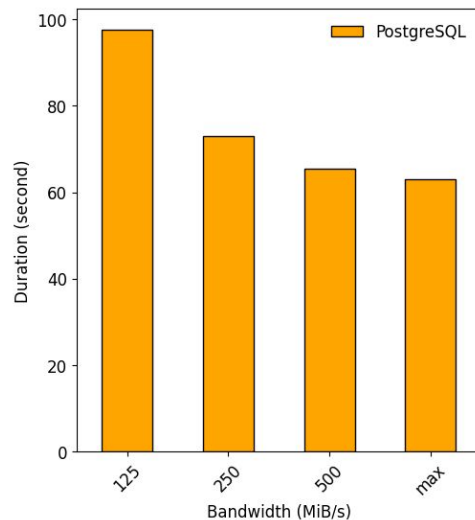
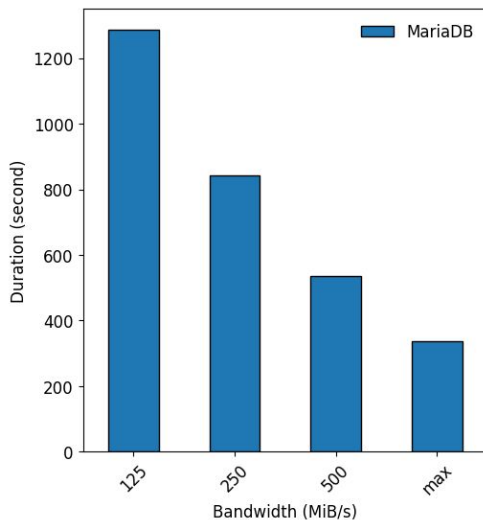
## Evaluation Environment

- Google Compute Engine with Ubuntu 22.04.2 LTS (GNU/Linux 6.1.14 x86\_64)
- Intel Xeon processor @ 2.20 GHz
- Processor has 24 cores and 128 GiB of memory
  - 12 cores and 112 GiB of memory dedicated to NVMeVirt
  - 12 cores and 16 GiB of memory dedicated to database engine
- configured with recommended configurations from optimization tools



# Evaluation

- OLAP workload with *TPC-H (Query 18 - Join & Aggregation)*



- For OLAP workload, the results we observed for OLTP workload didn't appear.
  - Hints us that *concurrency control* was what makes PostgreSQL more sensitive to I/O bandwidth for OLTP workload.

# Multi-Version Concurrency Control (MVCC)

- MVCC is the most popular transaction management scheme in modern DBMSs.
- With MVCC, the DBMS maintains multiple physical versions of a single logical object in the database.
  - Writers do not block readers and readers do not block writers.
- MVCC Design Decision:
  - Concurrency Control Protocol
  - **Version Storage**
  - **Garbage Collection**
  - **Index Management**

	Protocol	Version Storage	Garbage Collection	Indexes
<i>Postgres</i>	MV-2PL/SSI	Append-Only (O2N)	Tuple-level (VAC)	Physical
<i>MariaDB-InnoDB</i>	MV-2PL	Delta	Tuple-level (VAC)	Logical

<MVCC Implementations>

# Version Storage

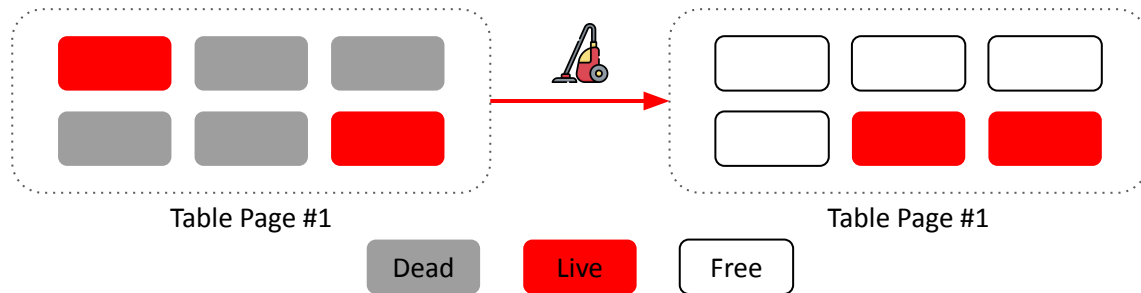
- DBMS always constructs a new physical version of a tuple when a transaction updates it.
- DBMS uses the tuples' pointer field to create a *version chain* per logical tuple.
  - This allows the DBMS to find the version that is visible to particular transaction at runtime.
  - Index pointers always point to the “head” of the chain.

## PostgreSQL: Append-Only Storage

### Main Table

	VALUE	POINTER
A <sub>0</sub>	\$111	●
A <sub>1</sub>	\$222	●
B <sub>1</sub>	\$10	∅
A <sub>2</sub>	\$333	∅

- Copies the whole content of the current version and store it in the same table space.
- Vacuum process periodically scan the table and look for reusable versions.



# Version Storage

- DBMS always constructs a new physical version of a tuple when a transaction updates it.
- DBMS uses the tuples' pointer field to create a *version chain* per logical tuple.
  - This allows the DBMS to find the version that is visible to particular transaction at runtime.
  - Index pointers always point to the “head” of the chain.

## MariaDB: Delta Storage

### Main Table

	VALUE	POINTER
A <sub>3</sub>	\$333	●
B <sub>1</sub>	\$10	

### Delta Storage Segment (= UNDO log)

	DELTA	POINTER
A <sub>1</sub>	(VALUE->\$111)	∅
A <sub>2</sub>	(VALUE->\$222)	●



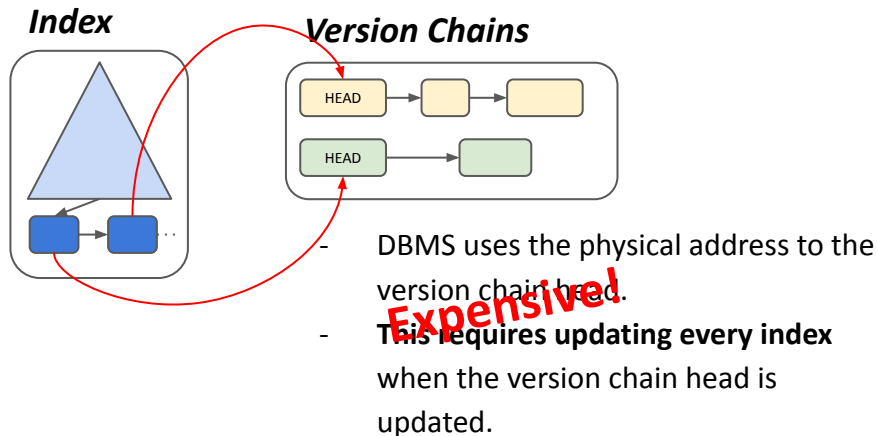
- InnoDB purge thread(s) will try free UNDO log pages containing old UNDO log records.

- On every update, *copy only the values that were modified to the delta storage and overwrite the master version.*

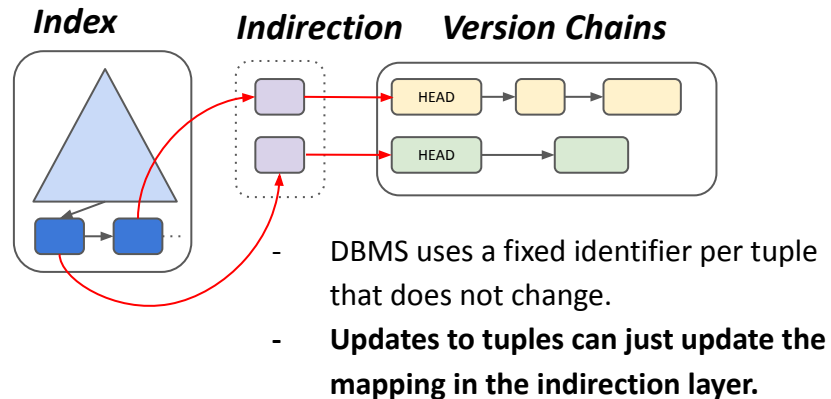
# Index Management

- Primary key indexes always point to version chain head.
- Using secondary index, tuple can be accessed using some piece of information other than the primary key.
  - PostgreSQL uses *physical pointers* to manage secondary indexes.
  - MariaDB uses *logical pointers* to manage secondary indexes.

## PostgreSQL: Physical Pointers



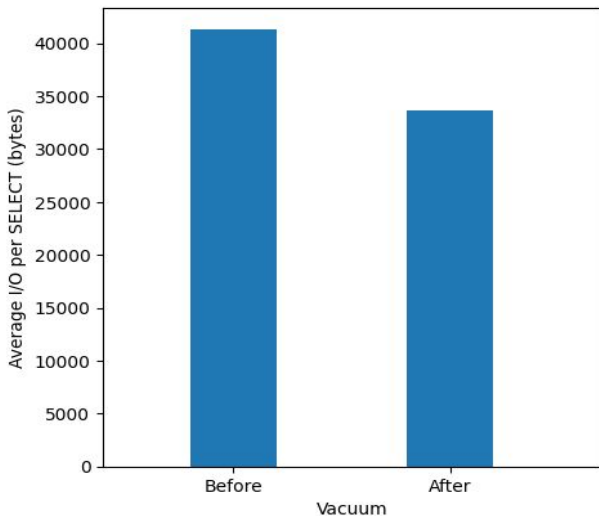
## MariaDB: Logical Pointers



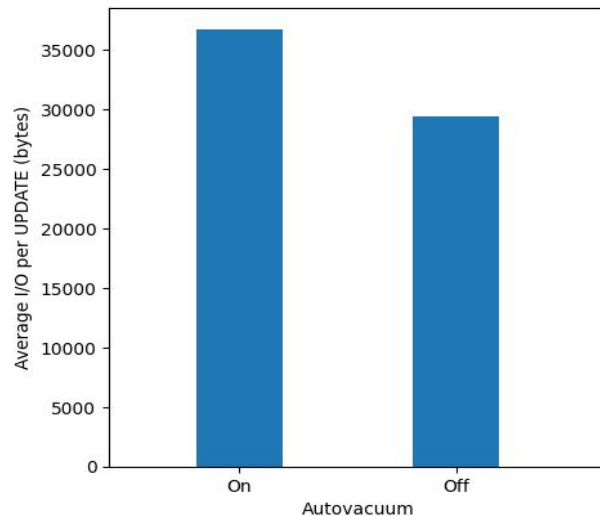


# Experimental Analysis

- Traversing a version chain to find the latest version requires more I/O.
- Inefficiency of *PostgreSQL*'s version storage method results in increased vacuum overhead related to I/O.



(a) Version chain scan overhead



(b) GC overhead

# Conclusion

- *PostgreSQL* is designed to utilize the storage device more eagerly than *MariaDB*.
  - Because of **the difference of their MVCC implementations!**
    1. Version Storage: append-only storage (*PostgreSQL*) vs. delta storage (*MariaDB*)
    2. Garbage Collection: background process vacuum (*PostgreSQL*) vs. delete undo log (*MariaDB*)
    3. Index Management: physical pointer (*PostgreSQL*) vs. logical pointer (*MariaDB*)
- As a result, these factors result in:
- Write Amplification
  - Expensive Garbage Collection
  - Occupying a large amount of bandwidth